

# Two Vignettes from Optimization

---


Qijia Jiang

April 2021

Stanford University



# Motivation (for Studying Optimization)

- Workhorse behind machine learning empirical success stories; applications across science & engineering domains – better understanding informs practice (e.g., hyper-parameter tuning)



# Motivation (for Studying Optimization)

- Workhorse behind machine learning empirical success stories; applications across science & engineering domains – better understanding informs practice (e.g., hyper-parameter tuning)
- Modern computational model adds new dimension to the traditional setup (distributed, communication, synchrony, robustness etc.) – invite exploration for different trade-offs



# Motivation (for Studying Optimization)

- Workhorse behind machine learning empirical success stories; applications across science & engineering domains – better understanding informs practice (e.g., hyper-parameter tuning)
- Modern computational model adds new dimension to the traditional setup (distributed, communication, synchrony, robustness etc.) – invite exploration for different trade-offs
- Tools can be brought to bear in a broader context – connections to MCMC sampling, online learning etc.

**Goal:** Complexity-theoretic investigations for various function classes under (non-classical) oracle models

- Higher Derivative Oracle  $\Rightarrow$  Highly smooth functions
- Parallel Oracle  $\Rightarrow$  Non-smooth functions

# Higher-Order Acceleration for Highly-Smooth Functions

---

# Smooth Function Minimization: Gradient Descent

## Setting

- Assumption: (1) Lipschitz gradient, i.e.,  
 $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$ ; (2)  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  is convex & differentiable
- Gradient Oracle: access to  $f(\cdot)$  and  $\nabla f(\cdot)$  at any query point  $x$
- Goal: find  $\epsilon$ -optimal point, i.e.,  $f(x_k) - f^* \leq \epsilon$  (unconstrained)

# Smooth Function Minimization: Gradient Descent

## Assumption

(1) Lipschitz gradient, i.e.,  $\|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$ ; (2)  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  is convex & differentiable.

## Gradient Descent

At each iteration  $k$ ,

$$\begin{aligned}x_{k+1} &= x_k - \frac{1}{L} \nabla f(x_k) \\ &= \arg \min_y \left\{ f(x_k) + \langle \nabla f(x_k), y - x_k \rangle + \frac{L}{2} \|y - x_k\|^2 \right\}\end{aligned}$$

Error decreases as  $\mathcal{O}(1/k)$ , dimension-free.



# Accelerated Gradient Descent

## Accelerated Gradient Descent [Nesterov '83, Zhu & Orecchia '17]

1. At each iteration  $k$ , compute

$$a_{k+1} = \frac{1}{2} \left( 1/L + \sqrt{1/L^2 + 4A_k/L} \right) \quad \text{and} \quad A_{k+1} = A_k + a_{k+1}$$

2. Set

$$\tilde{x}_k = \frac{A_k}{A_{k+1}} y_k + \frac{a_{k+1}}{A_{k+1}} x_k$$

3. Gradient Descent step:

$$y_{k+1} = \tilde{x}_k - \frac{1}{L} \nabla f(\tilde{x}_k) = \arg \min_y \left\{ \langle \nabla f(\tilde{x}_k), y - \tilde{x}_k \rangle + \frac{L}{2} \|y - \tilde{x}_k\|^2 \right\}$$

4. Mirror Descent step:

$$x_{k+1} = x_k - a_{k+1} \nabla f(\tilde{x}_k) = \arg \min_x \left\{ \langle a_{k+1} \nabla f(\tilde{x}_k), x - x_k \rangle + \frac{1}{2} \|x - x_k\|^2 \right\}$$

Error decrease as  $\mathcal{O}(\frac{1}{k^2})$ , although not a descent method.

# Lower Bound for First-Order Oracle Model

## Lower Bound [Nemirovski & Yudin '83]

Let  $\{x_k\}$  be the sequence of points generated by any first-order method satisfying  $x_k \in x_0 + \text{span}\{\nabla f(x_0), \dots, \nabla f(x_{k-1})\} \forall k \geq 1$  and assume  $f(\cdot)$  is  $L$ -smooth, convex, we have

$$f(x_k) - f^* \geq \Omega\left(\frac{L}{k^2}\right).$$

# Higher Order Oracle Model

## Setting

- Assumption: (1)  $f$  convex & differentiable; (2)  $p$ -th order smooth

$$\|\nabla^p f(x) - \nabla^p f(y)\| := \max_{\|v\|=1} |\nabla^p f(x)[v]^p - \nabla^p f(y)[v]^p| \leq L_p \|x - y\|$$

- $p$ -th Order Oracle: access to  $\{f(x), \nabla f(x), \dots, \nabla^p f(x)\}$
- Goal: find  $\epsilon$ -optimal point, i.e.,  $f(x_k) - f^* \leq \epsilon$
- Denote  $p$ -th order Taylor Expansion around  $x$  as

$$f_p(y; x) := f(x) + \sum_{i=1}^p \frac{1}{i!} \nabla^i f(x)[y - x]^i$$

# Higher Order Oracle Model

## Setting

- Assumption: (1)  $f$  convex & differentiable; (2)  $p$ -th order smooth
- $p$ -th Order Oracle: access to  $\{f(x), \nabla f(x), \dots, \nabla^p f(x)\}$
- Goal: find  $\epsilon$ -optimal point, i.e.,  $f(x_k) - f^* \leq \epsilon$

## Lower Bound [Agarwal & Hazan '18, Nesterov '18]

Let  $\{x_k\}$  be the sequence of points generated by some tensor method of degree  $p$  that satisfies mild assumption, have

$$\min_{0 \leq t \leq k} f(x_t) - f^* \geq \Omega \left( \frac{Lp}{k^{\frac{3p+1}{2}}} \right).$$

# Higher Order Oracle Model

## Setting

- Assumption: (1)  $f$  convex & differentiable; (2)  $p$ -th order smooth
- $p$ -th Order Oracle: access to  $\{f(x), \nabla f(x), \dots, \nabla^p f(x)\}$
- Goal: find  $\epsilon$ -optimal point, i.e.,  $f(x_k) - f^* \leq \epsilon$

### Lower Bound [Agarwal & Hazan '18, Nesterov '18]

Let  $\{x_k\}$  be the sequence of points generated by some tensor method of degree  $p$  that satisfies mild assumption, have

$$\min_{0 \leq t \leq k} f(x_t) - f^* \geq \Omega\left(\frac{L_p}{k^{\frac{3p+1}{2}}}\right).$$

### Accelerated Rate [Nesterov & Polyak '08, Baes '09, Nesterov '18]

There is an algorithm that achieves convergence rate  $\mathcal{O}\left(\frac{1}{k^{p+1}}\right)$ .

Coincide when  $p = 1$ . For  $p = 2$  : Accel Cubic-Regularized Newton.

# The Iteration-Complexity Optimal Algorithm

1. Compute  $\lambda_{k+1} > 0$  and  $y_{k+1} \in \mathbb{R}^d$  such that

$$y_{k+1} = \arg \min_y \left\{ f_p(y; \tilde{x}_k) + \frac{L_p}{p!} \|y - \tilde{x}_k\|^{p+1} \right\} \quad (1)$$

$$\frac{1}{2} \leq \lambda_{k+1} \frac{L_p \cdot \|y_{k+1} - \tilde{x}_k\|^{p-1}}{(p-1)!} \leq \frac{p}{p+1} \quad (2)$$

where

$$a_{k+1} = \frac{1}{2} \left( \lambda_{k+1} + \sqrt{\lambda_{k+1}^2 + 4\lambda_{k+1}A_k} \right), \quad A_{k+1} = A_k + a_{k+1}$$
$$\tilde{x}_k = \frac{A_k}{A_{k+1}} y_k + \frac{a_{k+1}}{A_{k+1}} x_k \quad (3)$$

2. Update

$$x_{k+1} = x_k - a_{k+1} \nabla f(y_{k+1})$$

## Convergence Guarantee [BJLLS, COLT'19]

Error decrease as  $\tilde{O}(k^{-\frac{3p+1}{2}})$ , i.e., after  $\tilde{O}(\epsilon^{-\frac{2}{3p+1}})$  calls to tensor minimization blackbox that solves (1), we find an  $\epsilon$ -optimal point.

# The Iteration-Complexity Optimal Algorithm

1. Compute  $\lambda_{k+1} > 0$  and  $y_{k+1} \in \mathbb{R}^d$  such that

$$y_{k+1} = \arg \min_y \left\{ f_p(y; \tilde{x}_k) + \frac{L_p}{p!} \|y - \tilde{x}_k\|^{p+1} \right\} \quad (1)$$

$$\frac{1}{2} \leq \lambda_{k+1} \frac{L_p \cdot \|y_{k+1} - \tilde{x}_k\|^{p-1}}{(p-1)!} \leq \frac{p}{p+1} \quad (2)$$

where

$$a_{k+1} = \frac{1}{2} \left( \lambda_{k+1} + \sqrt{\lambda_{k+1}^2 + 4\lambda_{k+1}A_k} \right), \quad A_{k+1} = A_k + a_{k+1}$$
$$\tilde{x}_k = \frac{A_k}{A_{k+1}} y_k + \frac{a_{k+1}}{A_{k+1}} x_k \quad (3)$$

## Remark 1:

Approximate joint optimization over “interpolation parameter” and next query point  $y_{k+1}$ , i.e.,

$$\text{prox}_f^p(y_k, x_k - y_k) := \arg \min_{y \in \mathbb{R}^d, \tau \in \mathbb{R}} f(y) + C \cdot \|y - (y_k + \tau(x_k - y_k))\|^{p+1}$$

implemented by p-th order Taylor expansion + 1D binary search.

# The Iteration-Complexity Optimal Algorithm

1. Compute  $\lambda_{k+1} > 0$  and  $y_{k+1} \in \mathbb{R}^d$  such that

$$y_{k+1} = \arg \min_y \left\{ f_p(y; \tilde{x}_k) + \frac{L_p}{p!} \|y - \tilde{x}_k\|^{p+1} \right\} \quad (1)$$

$$\frac{1}{2} \leq \lambda_{k+1} \frac{L_p \cdot \|y_{k+1} - \tilde{x}_k\|^{p-1}}{(p-1)!} \leq \frac{p}{p+1} \quad (2)$$

where

$$a_{k+1} = \frac{1}{2} \left( \lambda_{k+1} + \sqrt{\lambda_{k+1}^2 + 4\lambda_{k+1}A_k} \right), \quad A_{k+1} = A_k + a_{k+1}$$
$$\tilde{x}_k = \frac{A_k}{A_{k+1}} y_k + \frac{a_{k+1}}{A_{k+1}} x_k \quad (3)$$

2. Update

$$x_{k+1} = x_k - a_{k+1} \nabla f(y_{k+1})$$

## Remark 2:

Proof based on refinement of estimate sequence technique.

[Monteiro & Svaiter '13]



# Highly Parallel Oracle for Nonsmooth Functions

---

# Nonsmooth Optimization – Classical Model

## Setting

- Assumption: (1)  $f$  convex; (2) 1-Lipschitz:  $|f(x) - f(y)| \leq \|x - y\|$
- First Order Oracle: query  $x \in \mathbb{R}^d \Rightarrow$  return  $f(x), \nabla f(x)$
- Goal: find  $\epsilon$ -optimal point, i.e.,  $f(x_k) - f^* \leq \epsilon$

Sequential:  
( $k=1$ )





Figure 1: Schematic for Sequential Setup

# Nonsmooth Optimization – Classical Model




## Subgradient Descent

At iteration  $k$ ,

$$x_{k+1} \leftarrow x_k - \eta \nabla f(x_k)$$

Output:  $\bar{x}_K = \frac{1}{K} \sum x_k$

$\mathcal{O}(\frac{1}{\epsilon^2})$  queries suffice



## Cutting Plane Methods

High-dimensional binary search  
(separation oracle implementable  
by subgradient oracle thanks to  
convexity)

$\mathcal{O}(d \log(\frac{1}{\epsilon}))$  queries suffice

# Parallel Oracle

Allowed to submit  $K$  queries in parallel [Nemirovski '94].




Figure 1: Schematic for Parallel Setup

**Depth** # queries to parallel oracle

**Work** total # gradients computed / functions evaluated

$$\text{work} = K \times \text{depth}$$

# Parallel Oracle

Allowed to submit  $K$  queries in parallel [Nemirovski '94].

Lower bound (for  $K = 1$ ):

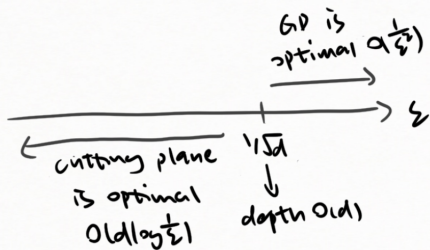


Figure 1: Upper & Lower Bound for Sequential

Question (for  $K > 1$ ):

For  $K = \text{poly}(d)$ , best possible depth? Power of adaptive information for convex optimization?

# Our Result

## Parallel Complexity [BJLLS, NeurIPS '19]

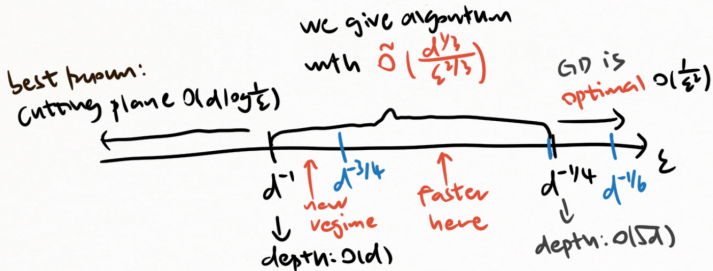


Figure 2: Upper & Lower Bound for Parallel

- $d^{-1/6}$  [Balanski & Singer '18]: GD is optimal for  $\epsilon = \tilde{\omega}(d^{-1/6})$
- $d^{-3/4}$  [Duchi, Bartlett, Wainwright '12]: Accelerated stochastic method with depth  $O(\frac{d^{1/4}}{\epsilon})$   
 → improve on sequential for  $\epsilon \in [d^{-3/4}, d^{-1/4}]$ , depth  $\in [\sqrt{d}, d]$

# Our Result

## How?

1. Convolve the nonsmooth function  $f$  with Gaussian kernel  $\gamma_r$ , i.e.,  $g(y) = \int f(y - x)\gamma_r(x) dx$ . Tradeoff between approximation vs. smoothness

$$|g(y) - f(y)| \leq r\sqrt{d} \quad \text{and} \quad \|\nabla^2 g(y)\|_{\text{op}} \leq \frac{1}{r}$$

2. Build inexact gradient oracle via sampling<sup>1</sup>, use to minimize the “prox” step approximately in depth=1

$$\tilde{\nabla} g(y) = \sum_{x_i} \gamma_r(y - x_i) \nabla f(x_i) \quad \text{and} \quad \min_y g(y) + Q(\|y - c\|)$$

3. Apply highly smooth acceleration result on the smoothed  $g(\cdot)$

<sup>1</sup>K ends up being  $\tilde{O}(d/\epsilon^2)$

## A More General Framework Emerges

Extend to more general “prox” function than  $\|\cdot\|^p$  – access to a subroutine that can approximately minimize

$$\arg \min_y f(y) + Q(\|x - y\|)$$

gives a systematic way of getting accelerated rate for the fcn class.

- Approximate Proximal Point:  $\approx \arg \min_y f(y) + L/2\|x - y\|^2$
- $p$ -th order Acceleration:  $= \arg \min_y f_p(y; x) + L_p/p!\|x - y\|^{p+1}$
- Highly Parallel:  $\approx \arg \min_y g(y) + C\|x - y\|^{p+1}$
- Ball-oracle [NeurIPS '20]:  $\approx \arg \min_{y: \|x-y\| \leq r} f(y)$   
efficiently implementable by trust-region for Hessian-stable functions, i.e.,  $c^{-1}\nabla^2 f(y) \preceq \nabla^2 f(x) \preceq c\nabla^2 f(y)$  for all  $\|x - y\| \leq r$



## A More General Framework Emerges

Extend to more general “prox” function than  $\|\cdot\|^p$  – access to a subroutine that can approximately minimize

$$\arg \min_y f(y) + Q(\|x - y\|)$$

gives a systematic way of getting accelerated rate for the fcn class.

- Approximate Proximal Point:  $\approx \arg \min_y f(y) + L/2\|x - y\|^2$
- $p$ -th order Acceleration:  $= \arg \min_y f_p(y; x) + L_p/p!\|x - y\|^{p+1}$
- Highly Parallel:  $\approx \arg \min_y g(y) + C\|x - y\|^{p+1}$
- Ball-oracle [NeurIPS '20]:  $\approx \arg \min_{y: \|x - y\| \leq r} f(y)$   
efficiently implementable by trust-region for Hessian-stable functions, i.e.,  $c^{-1}\nabla^2 f(y) \preceq \nabla^2 f(x) \preceq c\nabla^2 f(y)$  for all  $\|x - y\| \leq r$

Deeper implication for complexity theory – Better iteration complexity implemented by lower-order method.

Thanks!